

DEMO 10 - ACCOUNTING FOR MISSING DATA

ANSWER 1 - HANDLING MISSING DATA

1.0 - Prerequisite Check for Missing Data

The dataset given to us did not have predictor names but on this linke those attributes were available -

<http://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+%28Original%29>

Hence, after updating the

	ID	Clump_Thickness	Uniform_Cell_Size	Uniform_Cell_Shape	Marg_Adhesion
24	1057013	8	4	5	1
41	1096800	6	6	6	9
140	1183246	1	1	1	1
146	1184840	1	1	3	1
159	1193683	1	1	2	1
165	1197510	5	1	1	1
236	1241232	3	1	4	1
250	169356	3	1	1	1
276	432809	3	1	3	1
293	563649	8	8	8	1
295	606140	1	1	1	1
298	61634	5	4	3	1
316	704168	4	6	5	6
322	733639	3	1	1	1
412	1238464	1	1	1	1
618	1057067	1	1	1	1

	Single_Epith_Cell_Size	Bare_Nuclei	Bland_Chromatin	Normal_Nucleoli	Mitoses
24	2	NA	7	3	1
41	6	NA	7	8	1
140	1	NA	2	1	1
146	2	NA	2	1	1
159	3	NA	1	1	1
165	2	NA	3	1	1
236	2	NA	3	1	1
250	2	NA	3	1	1
276	2	NA	2	1	1
293	2	NA	6	10	1
295	2	NA	2	1	1
298	2	NA	2	3	1
316	7	NA	4	9	1
322	2	NA	3	1	1
412	1	NA	2	1	1
618	1	NA	1	1	1

	Class
24	1
41	0
140	0
146	0
159	0
165	0
236	0
250	0
276	0
293	1
295	0
298	0
316	0
322	0
412	0
618	0

Percentage of missing data: 2.289%

Please notice that 'Bare Nuclei' had all values assigned to NA which is below 5% threshold and justifies imputation.

In []: #CODE 1.0

```
require(ggthemes)
library(tidyverse)
library(magrittr)
```

```

library(TTR)
library(tidyr)
library(dplyr)
library(ggplot2)
library(plotly)
library(fpp2)
library(caTools)
library(reshape2)
library(psych)
require(graphics)
library(fBasics)
library(caret)
library(gridExtra)
library(DAAG)
library(rpart)
library(randomForest)
library(data.table)
library(mice)
library(MASS)
library(kknn)

# Read the file and name the columns
df<-read.table("breast-cancer-wisconsin.data.txt", stringsAsFactor = FALSE, header = F, sep = ",", na.strings="?")
colnames(df) <- c("ID", "Clump_Thickness", "Uniform_Cell_Size", "Uniform_Cell_Shape",
                 "Marg_Adhesion", "Single_Epith_Cell_Size", "Bare_Nuclei", "Bland_Chromatin",
                 "Normal_Nucleoli", "Mitoses", "Class")
df$Class <- as.factor(df$Class)
levels(df$Class) <- c(0, 1)

#which column contains the missing data
df[is.na(df$Bare_Nuclei),]
#check for % of missing observation (threshold < 5%)
print(sprintf("Percent of missing observation = %0.3f", 16/nrow(df)*100))

```

1.1 - Mean Imputation

The Bare Nuclei column is updated with mean values and first 24 values are displayed -

```

1.000000 10.000000 2.000000 4.000000 1.000000 10.000000 10.000000 1.000000
1.000000 1.000000 1.000000 1.000000 3.000000 3.000000 9.000000 1.000000
1.000000 1.000000 10.000000 1.000000 10.000000 7.000000 1.000000 3.544656

```

To do a sanity check, we also take a look at the mean of the column -

Average Value: 3.544656

```

In [ ]: #CODE - 1.1

#mean imputation
df.mean<-df
df.mean<-df.mean %>% mutate_at(vars(Bare_Nuclei),~ifelse(is.na(.x), mean(.x, na.rm = TRUE), .x))
#check it was imputed correctly
head(df.mean$Bare_Nuclei,24)

#check if mean was calculated correctly
mean(df.mean$Bare_Nuclei)

```

1.2 - Mode Imputation

The Bare Nuclei column is updated with mode values and first 24 values are displayed -

```

1 10 2 4 1 10 10 1 1 1 1 1 3 3 9 1 1 1 10 1 10 7 1 1

```

It doesn't tell us very much but the mode value of this column is - 1.

```

In [ ]: #CODE - 1.2

#found this mode function in the internet
getmode <- function(v) {
  uniqv <- unique(v)
  uniqv[which.max(tabulate(match(v, uniqv)))]
}
df.mode<-df
mode.result <- getmode(df.mode$Bare_Nuclei)

#fill NA with Mode of 1s
df.mode$Bare_Nuclei[is.na(df.mode$Bare_Nuclei)] <- mode.result
#check it was imputed correctly
head(df.mode$Bare_Nuclei,24)
print(mode.result)

```

1.3 - Regression Imputation

1.3 - Regression imputation

1 - First, we take a full model regression prediction of Bare Nuclei using all the other predictors -

```
Call:
lm(formula = Bare_Nuclei ~ Clump_Thickness + Uniform_Cell_Size +
    Uniform_Cell_Shape + Marg_Adhesion + Single_Epith_Cell_Size +
    Bland_Chromatin + Normal_Nucleoli + Mitoses, data = newdata.1)
```

Residuals:

Min	1Q	Median	3Q	Max
-9.7316	-0.9426	-0.3002	0.6725	8.6998

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-0.616652	0.194975	-3.163	0.00163	**
Clump_Thickness	0.230156	0.041691	5.521	4.83e-08	***
Uniform_Cell_Size	-0.067980	0.076170	-0.892	0.37246	
Uniform_Cell_Shape	0.340442	0.073420	4.637	4.25e-06	***
Marg_Adhesion	0.339705	0.045919	7.398	4.13e-13	***
Single_Epith_Cell_Size	0.090392	0.062541	1.445	0.14883	
Bland_Chromatin	0.320577	0.059047	5.429	7.91e-08	***
Normal_Nucleoli	0.007293	0.044486	0.164	0.86983	
Mitoses	-0.075230	0.059331	-1.268	0.20524	

Residual standard error: 2.274 on 674 degrees of freedom
Multiple R-squared: 0.615, Adjusted R-squared: 0.6104
F-statistic: 134.6 on 8 and 674 DF, p-value: < 2.2e-16

2 - Then we perform stepwise variable selection to get -

```
Call:
lm(formula = Bare_Nuclei ~ Clump_Thickness + Uniform_Cell_Shape +
    Marg_Adhesion + Bland_Chromatin, data = newdata.1)
```

Residuals:

Min	1Q	Median	3Q	Max
-9.8115	-0.9531	-0.3111	0.6678	8.6889

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-0.53601	0.17514	-3.060	0.0023	**
Clump_Thickness	0.22617	0.04121	5.488	5.75e-08	***
Uniform_Cell_Shape	0.31729	0.05086	6.239	7.76e-10	***
Marg_Adhesion	0.33227	0.04431	7.499	2.03e-13	***
Bland_Chromatin	0.32378	0.05606	5.775	1.17e-08	***

Residual standard error: 2.274 on 678 degrees of freedom
Multiple R-squared: 0.6129, Adjusted R-squared: 0.6107
F-statistic: 268.4 on 4 and 678 DF, p-value: < 2.2e-16

3 - Due to insignificant predictors we used stepwise forwards and backwards mode to build a more accurate model for Bare Nuclei.

	nvmax	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
1	1	2.548924	0.5228085	1.765516	0.2208983	0.08742394	0.1577356
2	2	2.432898	0.5624336	1.642851	0.2839852	0.10115619	0.1857652
3	3	2.362419	0.5900746	1.565549	0.2804584	0.09463731	0.1827845
4	4	2.278984	0.6185545	1.534310	0.2734543	0.08969728	0.1880005

Best model - nvmax : 4

After cross-validation, we infer that the model with these predictors -

1. Clump_Thickness
2. Uniform_Cell_Shape
3. Marg_Adhesion
4. Bland_Chromatin

have the **lowest RMSE value of 2.278** and **highest R-Square of 0.618!**

4 - After applying the new regression model, the Bare Nuclei column is updated with regressed values and first 24 values are displayed -

1 10 2 4 1 10 10 1 1 1 1 3 3 9 1 1 1 10 1 10 7 1 5

```
In [ ]: #CODE - 1.3

set.seed(123)

newdata<-df
missing.index<-which(is.na(newdata$Bare_Nuclei), arr.ind=TRUE)
newdata.1 <- newdata[-missing.index,2:10]
#all other predictors data points except for the missing value and response variable

#Linear Model
model <- lm(Bare_Nuclei~Clump_Thickness+
            Uniform_Cell_Size+
            Uniform_Cell_Shape+
            Marg_Adhesion+
            Single_Epith_Cell_Size+
            Bland_Chromatin+
            Normal_Nucleoli+
            Mitoses,data=newdata.1 )
summary(model)

# Fit the full model
full.model <- model
# Stepwise regression model
step.model <- stepAIC(full.model, direction = "both", trace = FALSE)
summary(step.model)

# Set seed for reproducibility
set.seed(123)
# Set up repeated k-fold cross-validation
train.control <- trainControl(method = "cv", number = 10)
# Train the model
step.model <- train(Bare_Nuclei ~., data = newdata.1 ,
                    method = "leapBackward",
                    tuneGrid = data.frame(nvmax = 1:4),
                    trControl = train.control
                    )

step.model$results
predicted.missing<-predict(step.model,newdata=df[missing.index,])
df.final.regression<-df
df.final.regression[missing.index,]$Bare_Nuclei<-as.integer(predicted.missing)#make predicted values integers
#final data with imputed regressed values
head(df.final.regression$Bare_Nuclei,24)
```

1.4 - Regression and Perturbation Imputation

1 - This warrants for the regression using random values for missing data, and this is final Bare Nuclei values plugged in -

1 10 2 4 1 10 10 1 1 1 1 3 3 9 1 1 1 10 1 10 7 1 5

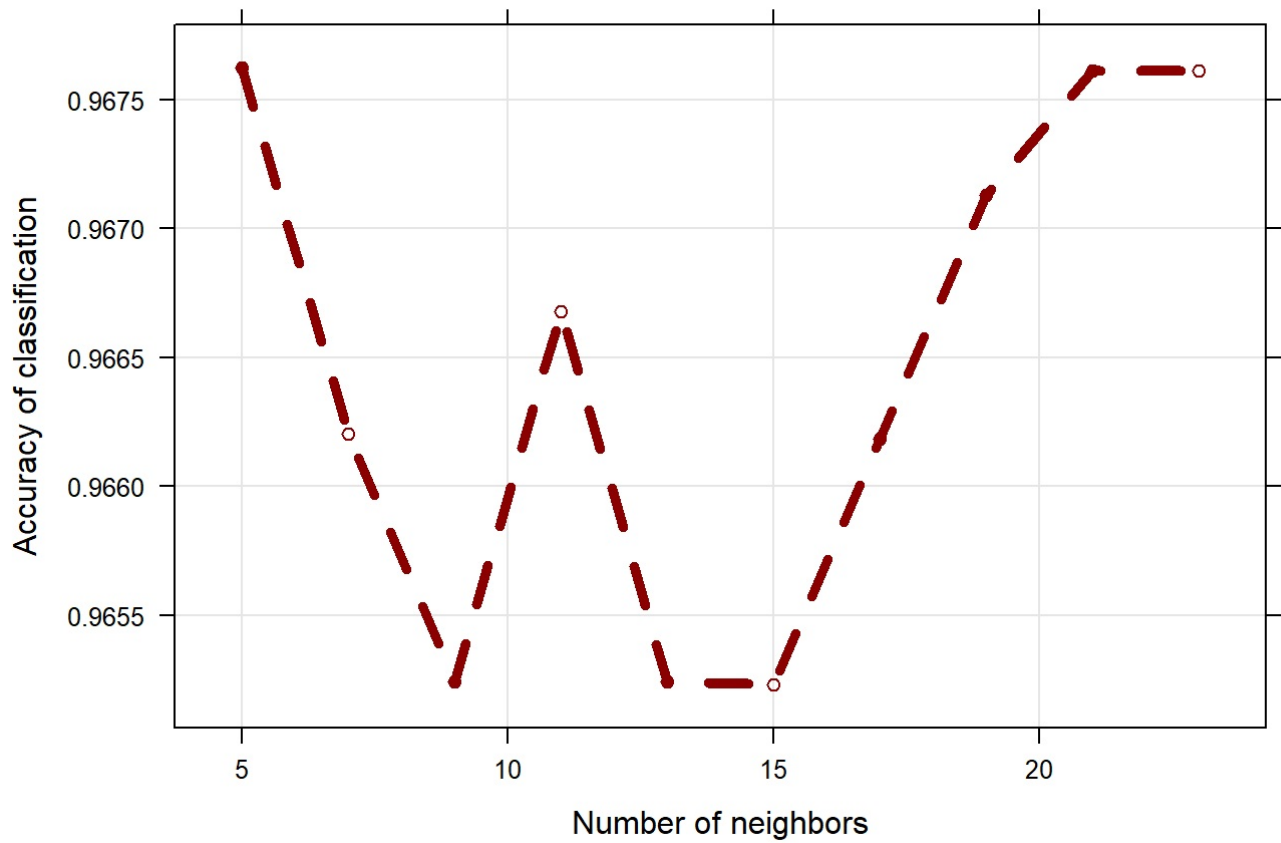
```
In [ ]: #CODE - 1.4

set.seed(123)
#generate 16 random numbers based off missing predicted values
n <- rnorm(16, mean = predicted.missing, sd = sd(predicted.missing))
#bounding the negative numbers to positive only
abs(n)
df.final.regression.pertubed<-df
df.final.regression.pertubed[missing.index,]$Bare_Nuclei<-as.integer(abs(n))#make predicted values integers
#final data with imputed perturbed regressed values
head(df.final.regression.pertubed$Bare_Nuclei,24)
```

1.5 - Comparisons

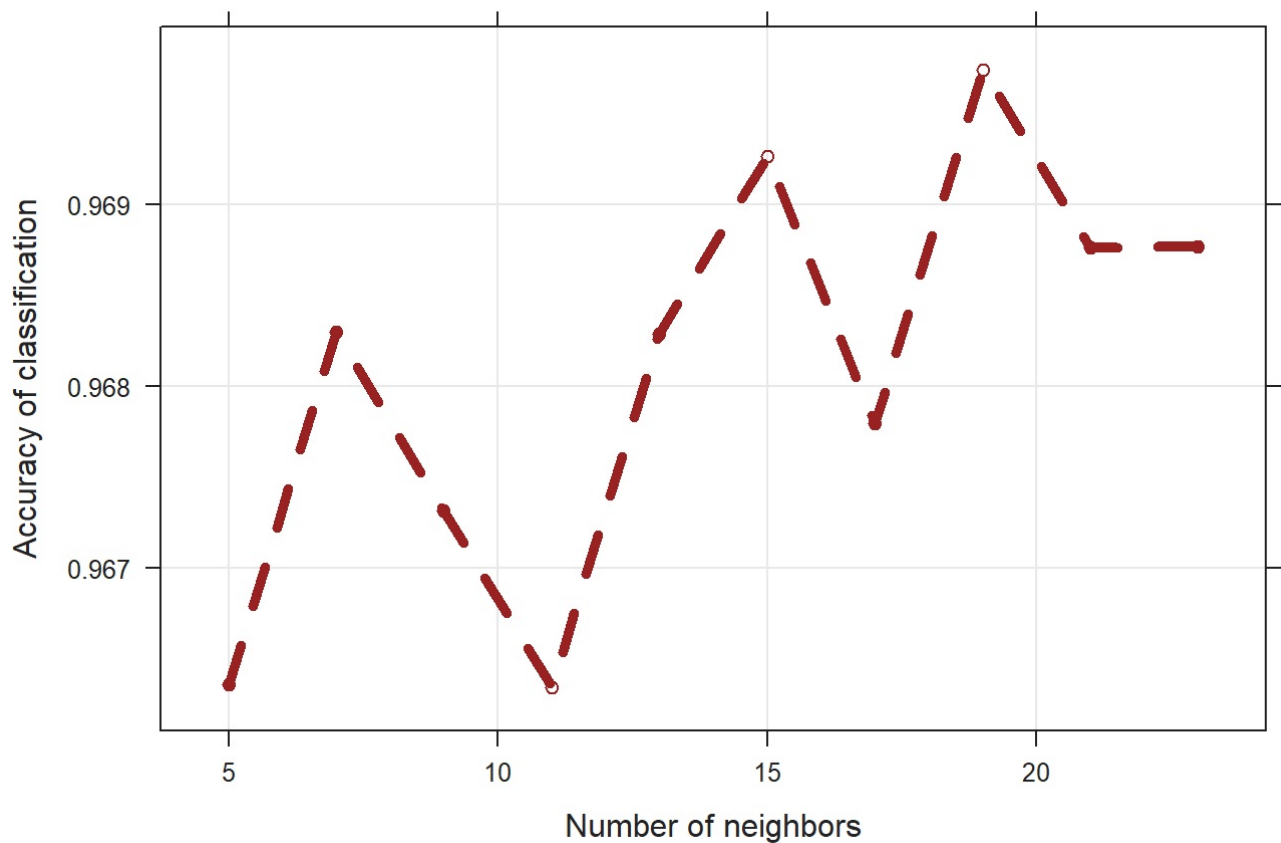
1 - KNN versus Mean Imputation

Mean Imputed Dataset: Accuracy of kNN with repeated 10-fold CV



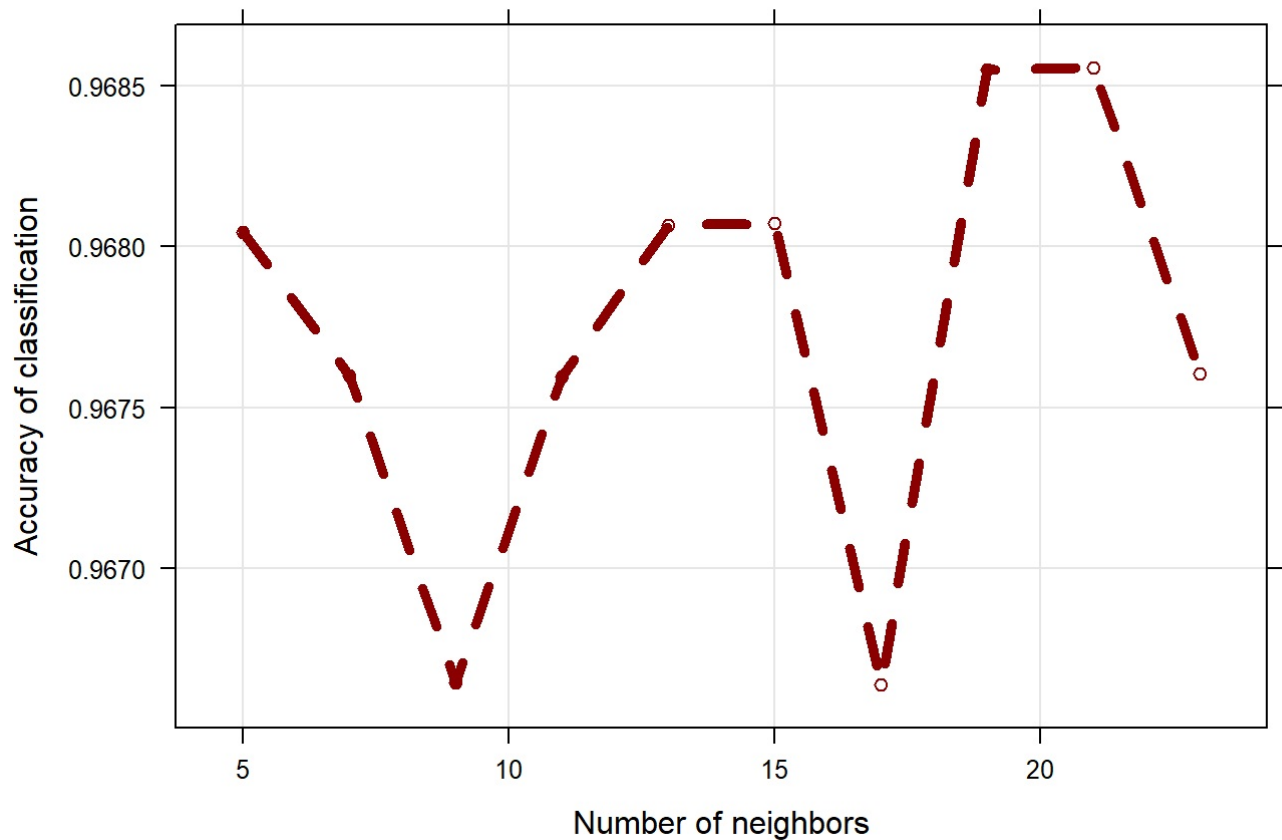
2 - KNN versus Removed Missing Points

Missing values removed: Accuracy of kNN with repeated 10-fold CV



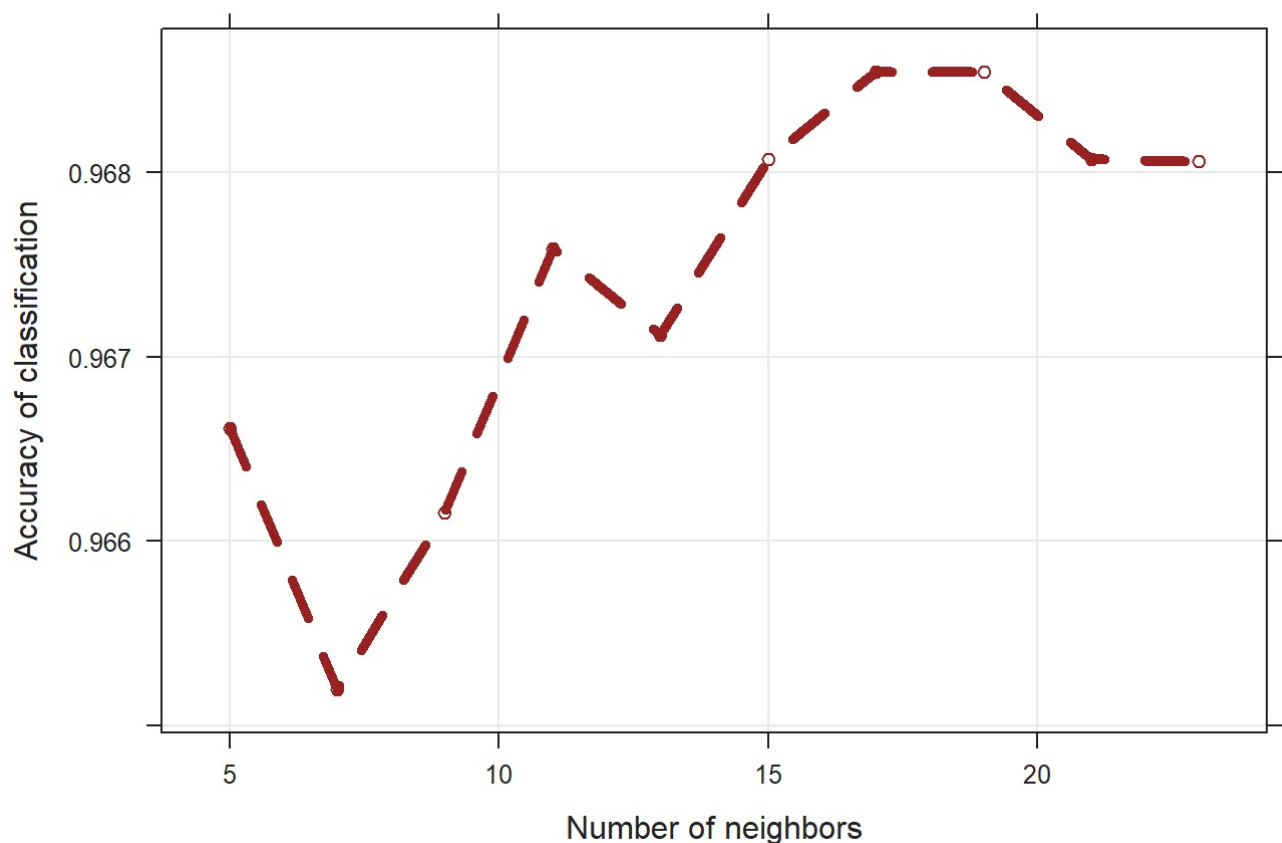
3 - KNN versus Regular Regression

Regression Imputed Dataset: Accuracy of kNN with repeated 10-fold CV



4 - KNN versus Perturbed Regression

Regression with Perturbation: Accuracy of kNN with repeated 10-fold CV



5 - Conclusion

With 10-fold CVV KNN accuracy was quite high - over 90% - but with differing k values across the techniques -

1. Mean approach K: 5

2. Missing approach K: 19
3. Perturbed regression approach K: 19
4. Regular regression approach K: 20

```
In [ ]: #CODE - 1.5

df1<-read.table("breast-cancer-wisconsin.data.txt",
               stringsAsFactor = FALSE, header = F, sep = ",", na.strings="?")
ctrl <- trainControl(method="repeatedcv",number=10,repeats = 3)
knn.mean <- train(x=df.mean[,1:10],y=as.factor(df.mean[,11]), method = "knn", trControl = ctrl,
                 preProcess = c("center","scale"), tuneLength = 10)

#Plot 1
plot(knn.mean,col = "dark red",lwd=5,lty=2,cex.lab=1.25,cex.main=1.5,
     main="Mean Imputed Dataset: Accuracy of kNN with repeated 10-fold CV",
     xlab="Number of neighbors",
     ylab="Accuracy of classification")

knn.regressed <- train(x=df.final.regression[,1:10],y=as.factor(df.final.regression[,11]), method = "knn",
                     trControl = ctrl,
                     preProcess = c("center","scale"), tuneLength = 10)

#plot 2
plot(knn.regressed,col = "dark red",lwd=5,lty=2,cex.lab=1.25,cex.main=1.5,
     main="Regression Imputed Dataset: Accuracy of kNN with repeated 10-fold CV",
     xlab="Number of neighbors",
     ylab="Accuracy of classification")

knn.perturbreg <- train(x=df.final.regression.perturbed[,1:10],
                      y=as.factor(df.final.regression.perturbed[,11]), method = "knn", trControl = ctrl,
                      preProcess = c("center","scale"), tuneLength = 10)

#Plot 3
plot(x=knn.perturbreg,col = "dark red",lwd=5,lty=2,cex.lab=1.25,cex.main=1.5,
     main="Regression with Perturbation: Accuracy of kNN with repeated 10-fold CV",
     xlab="Number of neighbors",
     ylab="Accuracy of classification")

df1.removed<-na.omit(df1)
knn.removed <- train(x=df1.removed[,1:10],y=as.factor(df1.removed[,11]), method = "knn", trControl = ctrl,
                   preProcess = c("center","scale"), tuneLength = 10)

#Plot 4
plot(knn.removed,col = "dark red",lwd=5,lty=2,cex.lab=1.25,cex.main=1.5,
     main="Missing values removed: Accuracy of kNN with repeated 10-fold CV",
     xlab="Number of neighbors",
     ylab="Accuracy of classification")
```

ANSWER 2 - REAL-LIFE OPTIMIZATION USE-CASE

On a past operational strategy endeavour at a non-profit in natural catastrophe management, a system of UAV vehicles needed to be set up along with their stationary charging banks. Then upon an incoming request for emergency delivery, a UAV would need to be selected and assigned to pickup commodity from source to destination and land at nearest charging bank to recharge to full fuel. The following describes the system's optimized work plan -

Decision Variables for drone 'n' upon service request 'r' -

1. Total Distance Covered = Distance from drone 'n' current location to commodity source location + Distance from commodity source location to delivery location + Distance between delivery location and nearest power charge + Distance covered in idle time to get back to original location + Deviation distance in between due to recharge
2. Fuel usage for drone 'n' for service request 'r' based on Total Distance Covered
3. Emergency level in measure of time of service request 'r'.

This is calculated across the entire system of drones and the constraints are as follows -

1. No service waits longer than time 't' if a drone is available.
2. Total fuel usage across 'k' drones having overlapping work schedules within period 'p' does not exceed $f(k,p)$.
3. Recharges are kept to a minimum of 'm' recharges across $f(k,p)$.

And for the objective function (IN THE FOLLOWING ORDER OF IMPORTANCE):

1. Minimize average delivery time.
2. Minimize total fuel usage.
3. Minimize recharge trips.

THE END-----

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js